

**CS-5340/6340, Programming Assignment #2**  
**Due: Sunday, October 25, 2015 by 11:00pm**

For this assignment, you should write a program that implements the Viterbi algorithm as applied to part-of-speech tagging. As we discussed in class, the Viterbi algorithm is an efficient way to find the most likely part-of-speech tag sequence for a sentence. As input, your program will need two files: (1) a file of probabilities, and (2) a file of sentences to process. Your program should read the names of the 2 input files from the command line. The command-line arguments should be in the following order:

```
viterbi <probabilities_file> <sentences_file>
```

For example, we should be able to run your program like so:

```
viterbi probs.txt sents.txt
```

## 1. The Probabilities File

This file will contain bigram and lexical generation probabilities. Each line will be formatted as:

X Y probability

which means that  $P(X | Y) = \text{probability}$ .

If both X and Y are part-of-speech tags, then it is a bigram probability. If X is a word and Y is a part-of-speech tag, then it is a lexical generation probability. To keep things simple for this assignment, you should assume that there are only 4 possible part-of-speech tags: *noun verb inf prep*, plus a special tag *phi* to denote the beginning of a sentence. (If you see any of these five strings, you can assume that they are tags and not words.) A sample probabilities file might look like this:

```
bears noun .02
bears verb .02
fish verb .07
fish noun .08
noun phi .80
verb phi .10
noun verb .77
etc.
```

Use a default value of .0001 for all bigram and lexical generation probabilities that are not present in the probabilities file. Your program should be case insensitive, so “bears”, “Bears”, and “BEARS” should all be treated as the same token.

## 2. The Sentences File

This file will contain a set of sentences that you should run your program on. Each sentence will be on a separate line. There will not be any punctuation marks in this file. A sample file might look like this:

```
This is a sentence
And here is another sentence
```

---

## OUTPUT FORMATTING

The output produced by your program should consist of 4 items for each sentence that is processed:

1. The sentence being processed.
2. The probabilities associated with **every** node in the Viterbi network when the algorithm is finished.
3. The values associated with **every** node in the Backpointer network (i.e., the data structure that you use to reconstruct the final tag sequence) when the algorithm is finished. NOTE: this information could be stored in the Viterbi network itself or in a separate data structure. Either way, please print these values separately from the Viterbi probabilities.
4. The probability and the part-of-speech tag assignments for the best tag sequence.

**IMPORTANT:** When you print this information, please format it exactly like the example in Figure 1 on the next page! Since you'll be printing a lot of numbers, this is crucial to ensure that we know exactly what your numbers and results correspond to.

When printing the probability values, please print **exactly 10 digits** to the right of the decimal point.

Please print the nodes as if you were reading the sentence left-to-right (so all nodes for the leftmost word in the sentence should appear first) and print the nodes within a column ordered as "noun", "verb", "inf", and "prep". The trace files that we give you will follow this convention.

The final part-of-speech tag assignments should be printed starting with the last word in the sentence and ending with the first word. This order should be the easiest one for you to produce as you traverse the backpointer network.

PROCESSING SENTENCE: bears fish

FINAL VITERBI NETWORK

$P(\text{bears}=\text{noun}) = 0.0160000000$

$P(\text{bears}=\text{verb}) = 0.0020000000$

$P(\text{bears}=\text{inf}) = 0.0000000100$

$P(\text{bears}=\text{prep}) = 0.0000000100$

$P(\text{fish}=\text{noun}) = 0.0001232000$

$P(\text{fish}=\text{verb}) = 0.0007280000$

$P(\text{fish}=\text{inf}) = 0.0000000440$

$P(\text{fish}=\text{prep}) = 0.0000004800$

FINAL BACKPTR NETWORK

Backptr(fish=noun) = verb

Backptr(fish=verb) = noun

Backptr(fish=inf) = verb

Backptr(fish=prep) = noun

BEST TAG SEQUENCE HAS PROBABILITY = 0.0007280000

fish -> verb

bears -> noun

Figure 1: Sample Output for Viterbi Program

## FOR CS-6340 STUDENTS ONLY! (30 pts)

In addition to the implementation of the Viterbi algorithm, CS-6340 students should also implement the Forward Algorithm to compute the cumulative evidence that each word has a specific part-of-speech.

In addition to the output of the Viterbi algorithm, your program should print the results of the Forward Algorithm by showing the probability computed for each node in the network. Please use the same printing conventions mentioned earlier regarding the order of the nodes and ten digits after the decimal point. For example, your output should be formatted like this:

### FORWARD ALGORITHM RESULTS

P(bears=noun) = 0.8888879012

P(bears=verb) = 0.1111109877

P(bears=inf) = 0.0000005556

P(bears=prep) = 0.0000005556

P(fish=noun) = 0.1447660127

P(fish=verb) = 0.8545600252

P(fish=inf) = 0.0000518360

P(fish=prep) = 0.0006221261

---

## GRADING CRITERIA

Your program will be graded based on new test cases! **So please test your program thoroughly to evaluate the generality and correctness of your code!** Even if your program works perfectly on the examples that we give you, that does not guarantee that it will work perfectly on additional test cases.

---

## ELECTRONIC SUBMISSION INSTRUCTIONS (a.k.a. “What to turn in and how to do it”)

You need to submit 3 things:

1. The source code files for your program. Be sure to include all files that we will need to compile and run your program!
2. A README file that includes the following information:
  - how to compile and run your code
  - which CADE machine you tested your program on  
(this info may be useful to us if we have trouble running your program)
  - any known bugs, problems, or limitations of your program

REMINDER: your program **must** compile and run on the linux-based CADE machines! We will not grade programs that cannot be run on the linux-based CADE machines.

3. Submit one trace file called **viterbi.trace** that shows the output of your program when given the input files on the CS-5340 web page.

You can generate a trace file in (at least) 3 different ways: (1) print your program's output to a file called *viterbi.trace*, (2) print your program's output to standard output and then pipe it to a file (e.g., *viterbi probs.txt sents.txt > viterbi.trace*), or (3) print your output to standard output and use the unix *script* command before running your program on the test files. The sequence of commands to use is:

```
script viterbi.trace
viterbi probs.txt sents.txt
exit
```

This will save everything that is printed to standard output during the session to a file called *viterbi.trace*.

---

To turn in your files, the CADE provides a web-based facility for electronic handin, which can be found here:

<https://webhandin.eng.utah.edu/>

Or you can log in to any of the CADE machines and issue the command:

```
handin cs5340 viterbi <filename>
```

---

**HELPFUL HINT:** you can get a listing of the files that you've already turned in via electronic handin by using the 'handin' command without giving it a filename. For example:

```
handin cs5340 viterbi
```

will list all of the files that you've turned in thus far. If you submit a new file with the same name as a previous file, the new file will overwrite the old one.